

Homework 1 answers

Note: For full credit, show your work!

You are welcome to discuss the problems with others, but write up your own solutions.

① For the matrix

$$A = \begin{pmatrix} 1 & 0 & 0 & 1/3 & 1/3 & 1/3 \\ 0 & 1 & 0 & 1/3 & 1/3 & 1/3 \\ 0 & 0 & 1 & 1/3 & 1/3 & 1/3 \\ 0 & 0 & 0 & 1/3 & 1/3 & 1/3 \\ 0 & 0 & 0 & 1/3 & 1/3 & 1/3 \\ 0 & 0 & 0 & 1/3 & 1/3 & 1/3 \end{pmatrix}$$

determine A^{300} .

First do the calculation in Matlab. Then do the same calculation by hand, justifying why the answer is correct.

Answer:

```
>> format rat
>> n = 3; A = [eye(n) ones(n,n)/n ; zeros(n,n) ones(n,n)/n]
```

A =

1	0	0	1/3	1/3	1/3
0	1	0	1/3	1/3	1/3
0	0	1	1/3	1/3	1/3
0	0	0	1/3	1/3	1/3
0	0	0	1/3	1/3	1/3
0	0	0	1/3	1/3	1/3

```
>> A^300
```

ans =

1	0	0	100	100	100
0	1	0	100	100	100
0	0	1	100	100	100
0	0	0	1/3	1/3	1/3
0	0	0	1/3	1/3	1/3
0	0	0	1/3	1/3	1/3

Why? A is a block matrix:

$$A = \begin{pmatrix} I & B \\ 0 & B \end{pmatrix} \quad \text{where } I \text{ is the } 3 \times 3 \text{ identity matrix}$$

$$B = \frac{1}{3} \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix}$$

$$\Rightarrow A^2 = \begin{pmatrix} I & I \cdot B + B^2 \\ 0 & B^2 \end{pmatrix} \quad \text{Since } B^2 = B, \quad A^2 = \begin{pmatrix} I & 2B \\ 0 & B \end{pmatrix}$$

Now let us prove by induction:

$$\text{Claim: For } k=1, 2, 3, \dots, \quad A^k = \begin{pmatrix} I & kB \\ 0 & B \end{pmatrix}$$

Proof:

Base case $k=1$: ✓

$$\text{Induction step: Assume } A^{k-1} = \begin{pmatrix} I & (k-1)B \\ 0 & B \end{pmatrix}$$

$$\begin{pmatrix} I & (k-1)B \\ 0 & B \end{pmatrix}$$

Induction step: Assume $A^{-1} = \begin{pmatrix} I & B \\ 0 & B^{-1} \end{pmatrix}$.

$$\begin{aligned} \Rightarrow A^k &= A \cdot A^{k-1} = \begin{pmatrix} I & B \\ 0 & B \end{pmatrix} \begin{pmatrix} I & (k-1)B \\ 0 & B \end{pmatrix} \\ &= \begin{pmatrix} I & (k-1)B + B^2 \\ 0 & B^2 \end{pmatrix} = \begin{pmatrix} I & kB \\ 0 & B \end{pmatrix} \checkmark \square \end{aligned}$$

② a) Give a 2×2 matrix A that transforms
 $(1, 0)$ to (a, c)

and $(0, 1)$ to (b, d) .

(That is, $A \begin{pmatrix} 1 \\ 0 \end{pmatrix} = \begin{pmatrix} a \\ c \end{pmatrix}$ and $A \begin{pmatrix} 0 \\ 1 \end{pmatrix} = \begin{pmatrix} b \\ d \end{pmatrix}$.)

Assuming $ad \neq bc$, give a 2×2 matrix that transforms

(a, c) to $(1, 0)$

and (b, d) to $(0, 1)$.

b) Give a 3×3 matrix A that transforms

$(1, 1, 1)$ to $(1, 0, 0)$,

$(2, -1, 0)$ to $(0, 1, 0)$

and $(0, 3, 3)$ to $(0, 0, 1)$.

(Feel free to use a computer, if that helps.)

c) Give a 3×3 matrix B that transforms

$(1, 1, 1)$ to $(5, 10, -1)$

$(2, -1, 0)$ to $(15, 0, 0)$

$(0, 3, 3)$ to $(-1, -1, 1)$.

(Hint: Use part (b).)

a) $A = \begin{pmatrix} a & b \\ c & d \end{pmatrix}$

$$A^{-1} = \frac{1}{ad-bc} \begin{pmatrix} d & -b \\ -c & a \end{pmatrix}$$

b) $A^{-1} = \begin{pmatrix} 1 & 2 & 0 \\ 1 & -1 & 3 \\ 1 & 0 & 3 \end{pmatrix} \Rightarrow A = \begin{pmatrix} 1 & 2 & -2 \\ 0 & -1 & 1 \\ -1/3 & -4/3 & 1 \end{pmatrix}$

c) $\begin{pmatrix} 5 & 15 & -1 \\ 10 & 0 & -1 \\ -1 & 0 & 1 \end{pmatrix} \times \begin{pmatrix} 1 & 2 & -2 \\ 0 & -1 & 1 \\ -1/3 & -4/3 & 1 \end{pmatrix} = \begin{pmatrix} 16/3 & -13/3 & 4 \\ 31/3 & 62/3 & -21 \\ -4/3 & -8/3 & 3 \end{pmatrix}$

③ Sketch the three lines

③ Sketch the three lines

$$x + 2y = 2$$


$$x - y = 2$$

$$y = 1$$


Are the equations solvable?

What happens if the right-hand sides are all 0?

Is there any nonzero choice for the right-hand sides that allows the three lines to intersect at the same point?



Since the three lines don't intersect, the equations are not solvable.



Now there is a
unique solution
 $x = y = 0$

($\vec{0}$ solves any
homogeneous system
of equations)

④ Describe the intersection of the three planes

$$u + v + w + z = 6,$$

$$u + w + z = 2, \text{ and}$$

$$u + w = 2,$$

all in four-dimensional space. Is it a line or a point or an empty set?

Answer: Use Gaussian elimination to find their intersection:

$$\begin{array}{l} -1 \uparrow \\ -6 \end{array} \left(\begin{array}{cccc|c} 1 & 1 & 1 & 1 & 6 \\ 1 & 0 & 1 & 1 & 2 \\ 1 & 0 & 1 & 0 & 2 \end{array} \right) \rightarrow \left(\begin{array}{cccc|c} 0 & 1 & 0 & 0 & 4 \\ 1 & 0 & 1 & 1 & 2 \\ 0 & 0 & 0 & 1 & 0 \end{array} \right)$$

\Rightarrow the solution is

$$v = 4, z = 0$$

$$u = 2 - w, \text{ where } w \text{ is free}$$

This is a line

What is the intersection if the fourth plane $u = -1$ is included? Find a fourth equation that leaves us with no solution.

If $u = -1$ is added, the solution is a point:

$$(u, v, w, z) = (-1, 4, 3, 0)$$

$u + w = 10$ will give no solution, for example

⑤ How fast is Matlab?

A useful skill is to be able to estimate how long a calculation is going to take before starting it. (For

example, you can determine how large a problem you can solve without buying a new computer!)

a) Matrix multiplication:

Here is some crude Matlab/Octave code for timing the multiplication of two random 1000×1000 matrices:

```
% time to multiply two random 1000x1000 matrices
```

```
n = 1000;
```

```
A = randn(n, n);
```

```
B = randn(n, n);
```

```
starttime = cputime;
```

```
C = A * B;
```

```
endtime = cputime;
```

```
elapsedtime = endtime - starttime
```

```
% same code, but averaged over 10 trials
```

```
trials = 10;
```

```
n = 1000;
```

```
A = randn(n, n);
```

```
B = randn(n, n);
```

```
starttime = cputime;
```

```
for i = 1:trials
```

```
    C = A * B;
```

```
end
```

```
endtime = cputime;
```

```
averageelapsedtime = (endtime - starttime)/trials
```

Python

```
import time
```

```
import numpy as np
```

```
n = 2000
```

```
trials = 10
```

```
totaltime = 0
```

```
for _ in range(trials):
```

```
    A = np.random.randn(n, n)
```

```
    B = np.random.randn(n, n)
```

```
    start = time.perf_counter()
```

```
    C = A.dot(B) # or np.dot(A,B)
```

```
                # NOT A * B!
```

```
    totaltime += time.perf_counter() - start
```

```
averagetime = totaltime / trials
```

```
averagetime
```

← CPU time
in seconds

(to reduce noise)

Your problem: Estimate the scaling of the running time for matrix multiplication, as n increases. That is, if the running time is $\approx c \cdot n^\alpha$ for some exponent α and constant c , estimate values for α and c .

(Don't worry about being too precise, but try to get something reasonable. For example, if $\alpha = 3$, then running the above code with $n=2000$ should take $8 = 2^\alpha$ times as long. Show your work!)

Based on your estimates for c and α , for how large an n can your computer multiply two random $n \times n$ matrices in one day (24 hours)?

(Again, don't worry about being too precise, or about what

happens when your computer runs out of memory.)

Answer: There are lots of ways of doing this, to find c and α to fit observed running times. We'll talk more about fitting curves later.

One simple solution to find α is to simply take the ratio of the average CPU time for $n=1000$ and $n=2000$, i.e.,

$$\frac{t_{2000}}{t_{1000}} = \frac{c \cdot 2000^\alpha}{c \cdot 1000^\alpha} = 2^\alpha$$

$$\Rightarrow \alpha = \log_2\left(\frac{t_{2000}}{t_{1000}}\right)$$

$$c = \frac{t_{1000}}{1000^\alpha}$$

Since 24 hours = 86,400 seconds, solving

$$86400 = c \cdot n^\alpha$$

$$\Rightarrow n = \left(\frac{86400}{c}\right)^{1/\alpha}$$

My Matlab code:

```
n = 1000;
totaltime = 0;
trials = 10;
for i = 1:trials
    A = randn(n, n);
    B = randn(n, n);
    starttime = cputime;
    C = A * B;
    endtime = cputime;
    totaltime = totaltime + (endtime - starttime);
end
t1000 = totaltime / trials
```

```
n = 2000;
totaltime = 0;
trials = 10;
for i = 1:trials
    A = randn(n, n);
    B = randn(n, n);
    starttime = cputime;
    C = A * B;
    endtime = cputime;
    totaltime = totaltime + (endtime - starttime);
end
t2000 = totaltime / trials
```

```
alpha = log2(t2000/t1000)
c = t1000 / 1000^alpha
n24 = (86400 / c)^(1/alpha)
```

alpha =

2.9069

c =

2.3211e-10

n24 =

1.0292e+05 $\approx 100,000$

Your answers will vary!

b) Solving a system of linear equations:

b) Solving a system of linear equations:

Repeat part a, but for solving n random linear equations in n variables.

Answer: Just as above, heres my code:

```
n = 1000;
totaltime = 0;
trials = 10;
for i = 1:trials
    A = randn(n, n);
    b = randn(n, 1);
    starttime = cputime;
    x = A \ b;
    endtime = cputime;
    totaltime = totaltime + (endtime - starttime);
end
t1000 = totaltime / trials
```

```
n = 2000;
totaltime = 0;
trials = 10;
for i = 1:trials
    A = randn(n, n);
    b = randn(n, 1);
    starttime = cputime;
    x = A \ b;
    endtime = cputime;
    totaltime = totaltime + (endtime - starttime);
end
t2000 = totaltime / trials
```

```
alpha = log2(t2000/t1000)
c = t1000 / 1000^alpha
n24 = (86400 / c)^(1/alpha)
```

alpha =

2.4829

c =

2.9529e-09

n24 =

2.6516e+05 $\approx 265,000$

c) Solving a **sparse** system of linear equations:

Repeat part a, but for solving

$$A\vec{x} = \vec{b},$$

where \vec{b} is a random vector of length n , and A is an $n \times n$ matrix with 10n random nonzero entries in random positions.

This code might be helpful:

```
% code for generating a random n x n matrix with one random non-zero
% entry in a random position
n = 1000;
A = sparse(n, n); % create an all-0 sparse matrix
i = randi(n); % randi(n) returns a random integer from 1 to n
j = randi(n);
A(i, j) = randn(1, 1); % set the i,j entry of A to a random value
```

Answer: As before,

```

n = 1000;
numnonzeroentries = 10 * n;
totaltime = 0;
trials = 30;
for i = 1:trials
    A = sparse(n, n);
    for k = 1:numnonzeroentries
        i = randi(n); j = randi(n);
        A(i, j) = randn(1,1);
    end
    b = randn(n, 1);
    starttime = cputime;
    x = A \ b;
    endtime = cputime;
    totaltime = totaltime + (endtime - starttime);
end
t1000 = totaltime / trials

alpha = log2(t2000/t1000)
c = t1000 / 1000^alpha
n24 = (86400 / c)^(1/alpha)

```

```

n = 2000;
numnonzeroentries = 10 * n;
totaltime = 0;
trials = 30;
for i = 1:trials
    A = sparse(n, n);
    for k = 1:numnonzeroentries
        i = randi(n); j = randi(n);
        A(i, j) = randn(1,1);
    end
    b = randn(n, 1);
    starttime = cputime;
    x = A \ b;
    endtime = cputime;
    totaltime = totaltime + (endtime - starttime);
end
t2000 = totaltime / trials

alpha =
    2.6274

c =
    8.3090e-10

n24 =
    2.1627e+05 ≈ 220,000

```

You may have noticed that the running time is highly variable. Running for more trials would help a little bit, but probably a different problem would give a more useful timing result.

Python:

a

```

import time
import numpy as np

```

```

n = 2000
trials = 10

```

```

totaltime = 0
for _ in range(trials):
    A = np.random.randn(n, n)
    B = np.random.randn(n, n)
    start = time.perf_counter()
    C = A.dot(B) # or np.dot(A,B)
    # NOT A * B!
    totaltime += time.perf_counter() - start

```

```

averagetime = totaltime / trials
averagetime

```

c

```

import time
import numpy as np
import scipy.sparse
from scipy.sparse.linalg import spsolve, bicg

```

```

n = 500

```

b

```

import time
import numpy as np

```

```

n = 1000
trials = 10

```

```


totaltime = 0
for _ in range(trials):
    A = np.random.randn(n, n)
    x = np.random.randn(n)
    b = A.dot(x)
    start = time.perf_counter()
    np.linalg.solve(A, b)
    totaltime += time.perf_counter() - start

```

```

averagetime = totaltime / trials
averagetime

```

```

import time
import numpy as np
import scipy.sparse
from scipy.sparse.linalg import spsolve, bicg

n = 500

total = 0
trials = 10
for _ in range(trials):
    sparsity = 10
    rows = np.random.randint(n, size=sparsity*n)
    cols = np.random.randint(n, size=sparsity*n)
    data = np.random.randn(sparsity*n)
    A = scipy.sparse.csr_matrix((data, (rows,cols)), shape=(n,n))
    #A += 1000* scipy.sparse.eye(n)  ## uncomment this to improve the condition no
    #A = A.toarray()                ## uncomment this to use dense matrix

    x = np.random.randn(n,1)
    b = A.dot(x)
    start = time.perf_counter()
    c = bicg(A,b)
    end = time.perf_counter()
    total += end - start
total

```