

Lecture 1: Motivation for linear algebra (class)

Admin: Textbook, syllabus, homework, midterms, final, grading, office hours, ...

MOTIVATION FOR LINEAR ALGEBRA

Theory: Linear transformations are everywhere!

- ~ Signals: Fourier transform is linear
- ~ Physics: Quantum time evolution is linear
- ~ Calculus: Integration and differentiation are linear

Applications: countless...


Application: Solving differential equations

example:
Solve

$$\begin{cases} f'(x) = g(x) \\ f(0) = f_0 \end{cases} \quad \forall x \in [0, 1] \quad \xrightarrow{\text{discretize}}$$

$$\begin{cases} f'(\frac{j}{n}) = g(\frac{j}{n}) \\ y_0 = f_0 \end{cases}$$

$$f(x) = f_0 + \int_0^x g(s) ds$$



$$f'(x) = \lim_{\delta \rightarrow 0} \frac{f(x+\delta) - f(x)}{\delta} \approx \frac{f(x+\frac{1}{n}) - f(x)}{1/n}$$

$$\underline{f'(\frac{j}{n}) \approx n(y_{j+1} - y_j) = g(\frac{j}{n})} \quad j = 0, 1, \dots, n-1$$

$$\begin{aligned} f'(x) &= \lim_{\delta \rightarrow 0} \frac{f(x+\delta) - f(x)}{\delta} = g(x) \\ &\approx \frac{f(x+\frac{1}{n}) - f(x)}{1/n} \end{aligned}$$

$$\Rightarrow n(y_{j+1} - y_j) \approx g(\frac{j}{n}), \quad j = 0, \dots, n-1$$

Aside:

$$f'(x) \approx \frac{1}{1/n} (f(x+\frac{1}{n}) - f(x))$$

$$f'(x) \approx \frac{1}{2/n} (f(x+\frac{1}{n}) - f(x-\frac{1}{n})) \leftarrow \text{better discretization}$$

Example: ECMWF 10-day weather forecasts



9km horizontal,
137 vertical levels
↓
 10^9 grid points
(~100 vars at each point)

<https://www.ecmwf.int/en/about/media-centre/news/2016/new-forecast-model-cycle-brings-highest-ever-resolution>

Application: Solving linear equations

eg. $\begin{cases} 2x - y = 3 \\ -x + y = -2 \end{cases}$

Gaussian elimination

$$\begin{cases} a(2x - y) = 3a - 2b \\ +b(-x + y) \end{cases}$$

$a=b=1:$
 $x=1 \Rightarrow y=-1$

Matlab

<https://matlab.mathworks.com/>

```
>> A = [2 -1; -1 1]
```

```
A =
```

```
2    -1
-1    1
```

```
>> b = [3; -2]
```

```
b =
```

```
3
-2
```

```
>> A \ b
```

```
ans =
```

```
1.0000
-1.0000
```

Python

<https://colab.research.google.com/>

```
[1] import numpy as np
```

```
A = [[2,-1], [-1,1]]
```

```
b = [3,-2]
```

```
np.linalg.solve(A, b)
```

```
[> array([ 1., -1.])
```

- but most applications are for **large** systems

we need **fast, approximate** solutions

often we solve the same system **repeatedly**

eg. $f'(x) = g(x)$

$f'(x) = h(x)$

$\Rightarrow n(y_{j+1} - y_j) = g(\delta/h) \Rightarrow n(y_{j+1} - y_j) = h(\delta/h)$

same coefficients of y_0, \dots, y_n

- what if **#equations > #variables**?

eg. $\begin{cases} x = 1 \\ x = 2 \end{cases}$

eg. $\begin{cases} 2x - y = 3 \\ -x + y = -2 \end{cases}$

$\begin{cases} x + y = 1 \\ 2x + 2y = 0 \end{cases}$

eg. $\begin{cases} x = 1 \\ x = 2 \end{cases}$
no solution!

Matlab

```
>> A = [1; 1];
>> b = [1; 2];
>> A \ b
ans =
```

$A \times x = b$

Matlab automatically returns $\arg \min \|Ax - b\|$

It tries to get as close as possible to a solution. We'll see more examples below.

eg. $\begin{cases} 2x - y = 3 \\ -x + y = -2 \\ x + y = 4 \end{cases}$
no solution!

```
>> A = [2 -1; -1 1; 1 1];
>> b = [3; -2; 4];
>> A \ b
ans =
```


np.linalg.lstsq

```
2.4286
1.2857
```

$\begin{cases} 2x + 2y = 0 \end{cases}$


- what if $\# \text{equations} < \# \text{variables}$?

Example: $x + y = 4$



infinitely many solutions!

$x + y + z = 4$



Matlab

<https://matlab.mathworks.com/>

```
>> A = [1 1];
>> b = [4];
>> A \ b
ans =
```

4

0

Python

<https://colab.research.google.com/>

```
import numpy as np
```

```
A = [[1, 1]]
b = [4]
np.linalg.solve(A, b)
```

```
-----
LinAlgError                                Traceback (most
<ipython-input-4-bc599599071c> in <module>()
      3 A = [[1, 1]]
      4 b = [4]
----> 5 np.linalg.solve(A, b)

<_array_function__ internals> in solve(*args, **kwargs)

1 frames
/usr/local/lib/python3.6/dist-packages/numpy/linalg/linalg
211     m, n = a.shape[-2:]
212     if m != n:
--> 213         raise LinAlgError('Last 2 dimensions o
214
215 def _assert_finite(*arrays):
```

LinAlgError: Last 2 dimensions of the array must be square

Example: "Compressed Sensing"


```
import matplotlib.pyplot as plt
```

```
w = 32
```

```
I = phantom(n=w)
```

```
imgplot = plt.imshow(I, cmap='gray')
imgplot
```

```
<matplotlib.image.AxesImage at 0x7f4...
```



https://en.wikipedia.org/wiki/Shepp-Logan_phantom

Shepp-Logan phantom

From Wikipedia, the free encyclopedia

The Shepp-Logan phantom is a standard test image created by Larry Shepp and Benjamin F. Logan for their 1974 paper *The Fourier Reconstruction of a Head Section*.^[1] It serves as the model of a human head in the development and testing of image reconstruction algorithms.^{[2][3][4]}



```
import numpy as np

x = np.ndarray.flatten(I)

n = len(x)
m = int(.7 * n)
print(m, n)

A = np.random.rand(m, n)

b = A.dot(x)

716 1024
```

716 equations
on $32^2 = 1024$ variables

<https://numpy.org/doc/stable/reference/generated/numpy.linalg.lstsq.html>

numpy.linalg.lstsq

numpy.linalg.lstsq(a, b, rcond='warn')

[source]

Return the least-squares solution to a linear matrix equation.


Computes the vector x that approximately solves the equation $a @ x = b$. The equation may be under-, well-, or over-determined (i.e., the number of linearly independent rows of a can be less than, equal to, or greater than its number of linearly independent columns). If a is square and of full rank, then x (but for round-off error) is the "exact" solution of the equation. Else, x minimizes the Euclidean 2-norm $\|b - ax\|$.

example:

```
import numpy as np

A = [[1, 1]]
b = [[4]]
np.linalg.lstsq(A, b, rcond=None)[0]

array([[2.],
       [2.]])
```




```
y = np.linalg.lstsq(A, b, rcond=None)[0]
```

```
y = np.resize(y, (w,w))
```

```
plt.imshow(y, cmap='gray')
```

<matplotlib.image.AxesImage at 0x7f4c922i>



```
y = np.linalg.solve(A, b)

-----
LinAlgError                                Traceback (most recent call last)
<ipython-input-4-9087bf01a13e> in <module>()
----> 1 y = np.linalg.solve(A, b)

__array_function__ internals in solve(*args, **kwargs)
-----
1 frames
/usr/local/lib/python3.6/dist-packages/numpy/linalg/linalg.py
211     m, n = a.shape[-2:]
212     if m != n:
--> 213         raise LinAlgError('Last 2 dimensions of
214                             the array must be square')
215 def _assert_finite(*arrays):
LinAlgError: Last 2 dimensions of the array must be square
```

What is CVXPY?

<https://www.cvxpy.org/>

CVXPY is a Python-embedded modeling language for convex optimization problems. It automatically transforms the problem into standard form, calls a solver, and unpacks the results.

```
import cvxpy as cp
```

```
# Construct the problem
```

```
z = cp.Variable(n)
```

```
objective = cp.Minimize(cp.sum_squares(A*z - b))
```

```
constraints = [0 <= z] # can also try [0 <= z, z <= 1]
```

```
problem = cp.Problem(objective, constraints)
```


```
# Solve it
```

```
result = problem.solve()
```

```
# The optimal value for z is stored in z.value
```

```
plt.imshow(np.resize(z.value, (w,w)), cmap='gray')
```

<matplotlib.image.AxesImage at 0x7f4c82b9f710>



Moral: For some applications we can solve for n variables with $< n$ equations!

"Compressed Sensing"

https://en.wikipedia.org/wiki/Compressed_sensing

for sparse solutions
(in some basis)

Applications


- 4.1 Photography
- 4.2 Holography
- 4.3 Facial recognition
- 4.4 Magnetic resonance imaging
- 4.5 Network tomography
- 4.6 Shortwave-infrared cameras
- 4.7 Aperture synthesis in radio astronomy
- 4.8 Transmission electron microscopy

Magnetic resonance imaging

Compressed sensing has been used^{[36][37]} to shorten magnetic resonance imaging scanning sessions on conventional hardware.^{[38][39][40]} Reconstruction methods include

- ISTA
- FISTA
- SiSTA
- ePRESS^[41]
- EWISTA^[42]
- EWISTARS^[43] etc.

Compressed sensing addresses the issue of high scan time by enabling faster acquisition by measuring fewer Fourier coefficients. This produces a high-quality image with relatively lower scan time. Another application (also discussed ahead) is for CT reconstruction with fewer X-ray projections. Compressed sensing, in this case, removes the high spatial gradient parts – mainly, image noise and artifacts. This holds tremendous potential as one can obtain high-resolution CT images at low radiation doses (through lower current-mA settings).^[44]



of variables
 $n = 128^2 = 16384$

of observations
 $m = 4480 = 0.27n$
 15 minutes
 in Li magic

Themes

- Geometry
 - linear transformations A
 - hyperplanes
 - High dimensions
 - sparse matrices ✓
 - dimension reduction ✓
 - Systems of linear equations $A\vec{x} = \vec{b}$
 - Computer-assisted linear algebra
 - Matlab
 - python/numpy
- DIY: how they work


Application: Image compression

Original

next channel


Result

Keep the largest 10% of coefficients in the Hadamard basis



not sparse!

Hadamard basis



sparse (in H basis)

Theme:
Sparse matrices



$$f'(x) = g(x)$$

$$y_{j+1} - y_j = \frac{1}{n} g(\delta/n), \quad j = 0, \dots, n-1$$

only 2 nonzero coefficients per equation

$$0 \cdot y_1 + 0 \cdot y_2 + \dots - 1 \cdot y_j + 1 \cdot y_{j+1} + 0 \cdot y_{j+2} + \dots$$

from scipy import sparse

```
n = 1000
A = sparse.diags([1, -1], offsets=[1, 0], shape=(n, n+1))
```

print(A)

```
(0, 1)    1.0    (0, 0)   -1.0
(1, 2)    1.0    (1, 1)   -1.0
(2, 3)    1.0
(3, 4)    1.0
(4, 5)    1.0
(5, 6)    1.0
(6, 7)    1.0
```

Sparse representations use less memory (and are faster to use)
print(A.size, A_denserep.size)

2000 10100
2000 10⁶

A_denserep = A.toarray()


print(A_denserep)

```
[[-1.  1.  0. ...  0.  0.  0.]
 [ 0. -1.  1. ...  0.  0.  0.]
 [ 0.  0. -1. ...  0.  0.  0.]
 ...
 [ 0.  0.  0. ...  1.  0.  0.]
 [ 0.  0.  0. ... -1.  1.  0.]
 [ 0.  0.  0. ...  0. -1.  1.]]
```

import matplotlib.pyplot as plt

plt.imshow(A_denserep, cmap='gray')

<matplotlib.image.AxesImage at 0x7f1

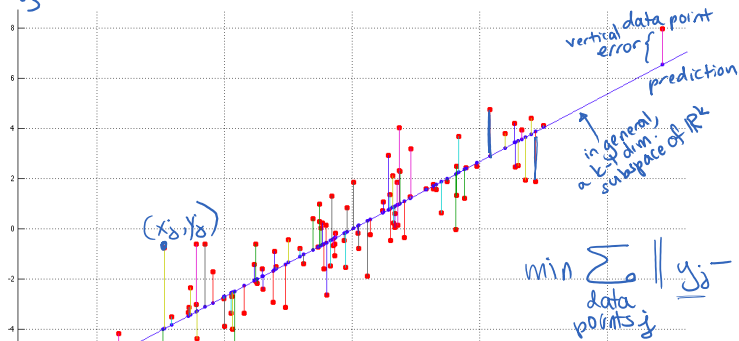


Theme:
Dimension reduction

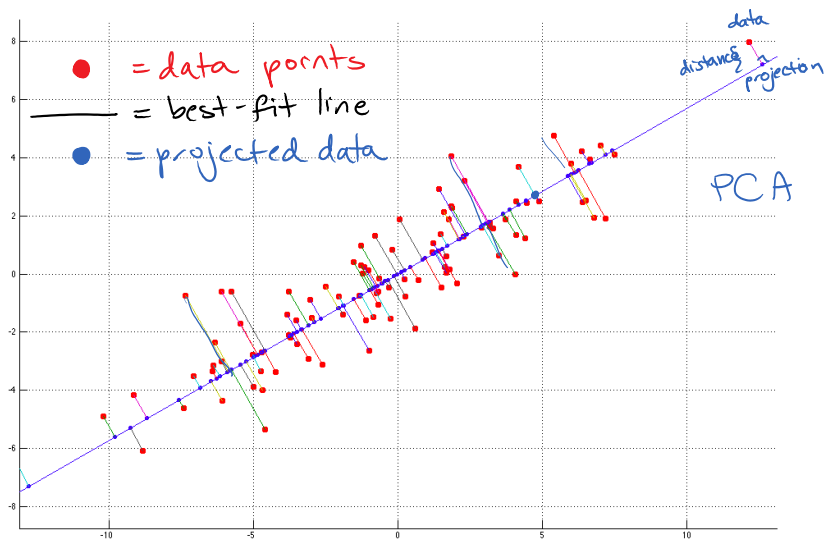
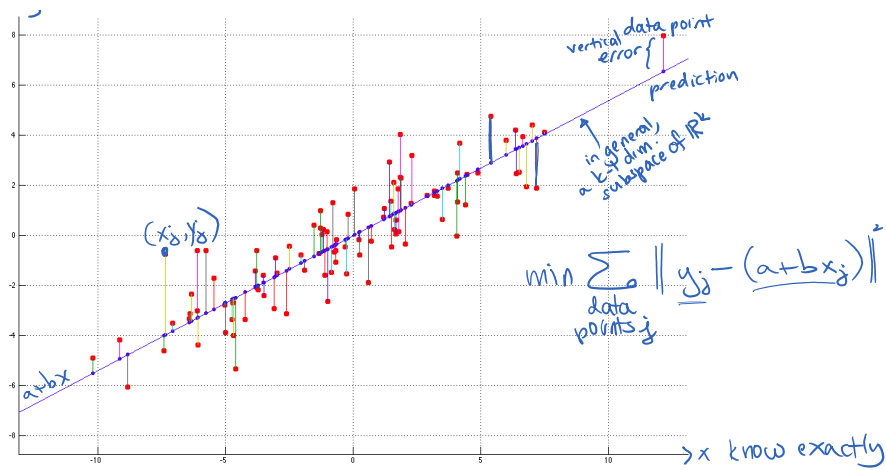
Applications: Least-squares fitting
Principal component analysis (PCA)

https://en.wikipedia.org/wiki/Principal_component_analysis

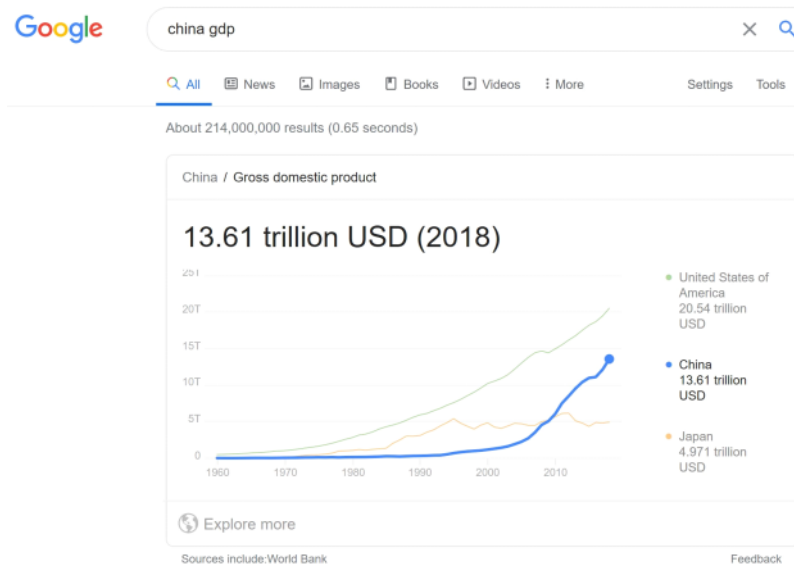
y noisy




$$\min \sum_{\text{data points } j} \|y_j - (a + bx_j)\|$$



Example: Predict China's gross domestic product (GDP) in 2030.
Answer:





$$G_t = G_0 e^{rt}$$

$$\Downarrow$$

$$\log G_t = (\log G_0) + rt$$

straight line

An exponential should fit the data better

Download the GDP data from the World Bank using pandas

```
import numpy as np
import pandas as pd


country = 'chn'
download_url = 'http://api.worldbank.org/v2/sources/2/country/' \
    + country + '/series/NY.GDP.MKTP.CD?format=json'
data = pd.read_json(download_url)
data = data['source']['data']

years = np.array([int(term['variable'][2]['value']) for term in data])
values = np.array([float(term['value']) for term in data])
```

Plot the data using matplotlib

```
import matplotlib.pyplot as plt
plt.style.use('seaborn-whitegrid')


plt.scatter(years, values, color='red')
plt.xlabel('Year')
plt.ylabel('GDP')
plt.show()
```




```
pred_years = np.arange(1970, 2031)
y = np.exp([fit.dot([1, year]) for year in pred_years])

plt.scatter(years, np.log(values), color='red')
plt.plot(pred_years, np.log(y), color='blue')
plt.xlabel('Year')
plt.ylabel('log(GDP)')
plt.show()


plt.scatter(years, values, color='red')
plt.plot(pred_years, y, color='blue')
plt.xlabel('Year')
plt.ylabel('GDP')
plt.show()
```



```
ListPlot[{#} & /@ dataprojectedmanually, AspectRatio -> 1, PlotMarkers -> data[[All, 1, 1]],  
PlotStyle -> (data[[All, 1, 3]]) /. {"D" -> Blue, "R" -> Red, "I" -> Black}]
```


```
{97, 17}
```

```
{17, 2}
```




- Even purely combinatorial applications

Spectral image segmentation



Spectral clustering

```
matrix = Import["/Users/breic/Desktop/adjacencymatrix.txt", "Table"];
matrix += Transpose[matrix];
matrix // ArrayPlot
```



```
laplacian = DiagonalMatrix[Plus@@# & /@ matrix] - matrix // N;
di = DiagonalMatrix[(1 / Plus@@#) & /@ matrix // N];
evs = Eigensystem[Sqrt[di].laplacian.Sqrt[di]] // Transpose // Reverse;
```

```
coordinates = evs[[2 ;; 9, 2]] // Transpose;
```

```
numclusters = 16; → embedding into  $\mathbb{R}^d$ 
```

```
ClusteringComponents[coordinates, numclusters, 1, Method -> "PAM"]
```

Indiana Jones and the L	A Walk to Remember	Bend It Like Beckham	Con Air
Lord of the Rings: The	Coyote Ugly	Bridget Jones's Diary	Double Jeopardy
Lord of the Rings: The	Dirty Dancing	Frida	Gone in 60 Seconds
Lord of the Rings: The	How to Lose a Guy in 10	Life Is Beautiful	Independence Day
Raiders of the Lost Ark	Maid in Manhattan	Love Actually	Lethal Weapon 4
Star Wars: Episode IV:	Pretty Woman	Moulin Rouge	Men in Black II
Star Wars: Episode VI:	Sister Act	My Big Fat Greek Weddin	Pearl Harbor
Star Wars: Episode V: T	The Princess Diaries 2:	Pride and Prejudice	The Fast and the Furiou
The Lord of the Rings:	The Princess Diaries (W	Rabbit-Proof Fence	The Patriot
	The Wedding Planner	Shakespeare in Love	Tomb Raider
	What Women Want	Whale Rider	Twister
12 Angry Men	A Bug's Life	Amelie	2001: A Space Odyssey
Airplane!	Breakfast at Tiffany's	American Beauty	All the President's Men
American Pie	City of Angels	Being John Malkovich	Blade Runner
American Pie 2	Ever After: A Cinderell	Crouching Tiger	Gandhi
Austin Powers in Goldme	Finding Nemo (Widescree	Election	Jaws
Austin Powers: Internat	Grease	Eternal Sunshine of the	L.A. Confidential
Austin Powers: The Spy	Harry Potter and the Ch	High Fidelity	Lawrence of Arabia
Interview with the Vamp	Harry Potter and the Pr	Lock	Lord of the Rings: The
Liar Liar	Harry Potter and the So	Lost in Translation	One Flew Over the Cucko
Meet the Parents	Runaway Bride	Magnolia	Seven Samurai
Ransom	The Lion King: Special	Run Lola Run	The Aviator
Spaceballs	The NeverEnding Story	Rushmore	The Exorcist
Spider-Man	The Princess Bride	Sideways	The Godfather
Wayne's World	The Sound of Music	The Royal Tenenbaums	The Godfather
	Willy Wonka & the Choco	Y Tu Mama Tambien	The Graduate
			The Great Escape
			The Maltese Falcon

Cold Mountain	A Fish Called Wanda	A Knight's Tale	Adaptation
Collateral	Alien: Collector's Edit	Ice Age	A Few Good Men
Crash	Back to the Future	Jurassic Park	Air Force One
Fahrenheit 9/11	Back to the Future Part	Lara Croft: Tomb Raider	Armageddon
Finding Neverland	Batman	Minority Report	Clear and Present Danger
Hotel Rwanda	Die Hard 2: Die Harder	Pirates of the Caribbean	Crimson Tide
Man on Fire	Die Hard With a Vengeance	Rush Hour	Enemy of the State
Master and Commander: T	Goldfinger	Rush Hour 2	Entrapment
Million Dollar Baby	Groundhog Day	Sleeping Beauty: Specia	High Crimes
Ocean's Twelve	Indiana Jones and the T	Spider-Man 2	In the Line of Fire
Ray	Men in Black	Star Wars: Episode II:	Lethal Weapon
Road to Perdition	Mission: Impossible	Star Wars: Episode I: T	Lethal Weapon 2
Runaway Jury	Predator: Collector's E	Terminator 3: Rise of t	Lethal Weapon 3
Seabiscuit	Rocky	The Fifth Element	Patriot Games
The Manchurian Candidat	Speed	The Incredibles	Rules of Engagement
The Notebook	Star Trek II: The Wrath	The Matrix	Swordfish
The Phantom of the Oper	Terminator 2: Extreme E	The Matrix: Reloaded	The Bone Collector
The Pianist	The Hunt for Red Octobe	The Matrix: Revolutions	The Client
	The Terminator	The Mummy	The Fugitive
	True Lies	The Mummy Returns	The Negotiator
		X2: X-Men United	The Pelican Brief
		X-Men	The Rock
			The Sum of All Fears
			Apollo 13
			As Good as It Gets
			Black Hawk Down
			Boys Don't Cry
			Cast Away
			Chocolat
12 Monkeys	Ace Ventura: Pet Detect	50 First Dates	Dances With Wolves: Spe
Almost Famous	A League of Their Own	Anger Management	Dead Man Walking
American History X	A River Runs Through It	Bad Boys II	Driving Miss Daisy
Anchorman: The Legend o	Basic Instinct	Behind Enemy Lines	Enemy at the Gates
Donnie Darko	Cheaper by the Dozen	Bruce Almighty	E.T. the Extra-Terrestre
Garden State	Daddy Day Care	Dodgeball: A True Under	Field of Dreams
GoodFellas: Special Edi	Erin Brockovich	Harold and Kumar Go to	Forrest Gump
Grosse Pointe Blank	Face/Off	Hero	Fried Green Tomatoes
Heat: Special Edition	Father of the Bride	Hidalgo	Gladiator
Kill Bill: Vol. 1	Kindergarten Cop	Hitch	Glory
Kill Bill: Vol. 2	Legally Blonde	Hostage	Good Will Hunting
Memento	Mrs. Doubtfire	I Robot	Jerry Maguire
Napoleon Dynamite	Notting Hill	Meet the Fockers	Moonstruck
Office Space	Pay It Forward	National Treasure	My Cousin Vinny
Pulp Fiction	Phenomenon	Ocean's Eleven	October Sky
Requiem for a Dream	Serendipity	Sahara	Philadelphia
Reservoir Dogs	Shall We Dance?	Shrek 2	Primal Fear
Seven	Sleepless in Seattle	The Bourne Identity	Rain Man
Sin City	Steel Magnolias	The Bourne Supremacy	Remember the Titans
		The Count of Monte Cris	

2. algorithms based on fast SDD solvers

* max-flow & multi-commodity flow problems

- for decades, best algorithms were deterministic, combinatorial
[Goldberg-Rao '98]: $\tilde{O}(m\sqrt{n}/\epsilon)$ augmenting paths
for $(1-\epsilon)$ -approx max flow blocking flows...

- recent breakthroughs based on numerical linear algebra,

spectral graph theory

[S-H Teng et al. '11]: $\tilde{O}(mn^{1/3}/\epsilon^{1/3})$
'usc

[Kelner et al., '13; Sherman '13]:

$\tilde{O}(m/\epsilon^2)$ or $\tilde{O}(km/\epsilon^2)$ for k flows

* generating random spanning trees

* graph sparsification

* sparsest cut

* distributed routing